

UNIVERSITY COLLEGE LONDON

DOCTORAL THESIS

---

**Song plays: Using machine learning to  
model the temporal patterns of music  
listeners**

---

*Author:*

**Badrul ALOM**

*Supervisor:*

Dr. Mark HERBSTER

*Advised by:*

Pedro Mediano (Emotech  
Ltd.)

*A thesis submitted in fulfillment of the requirements  
for the degree of Master of Science*

*in the*

Department of Computer Science

August 25, 2017



## Declaration of Authorship

I, Badrul ALOM, declare that this thesis titled, “Song plays: Using machine learning to model the temporal patterns of music listeners” and the work presented in it are my own. I confirm that:

- This work was done wholly or mainly while in candidature for a research degree at this University.
- Where any part of this thesis has previously been submitted for a degree or any other qualification at this University or any other institution, this has been clearly stated.
- Where I have consulted the published work of others, this is always clearly attributed.
- Where I have quoted from the work of others, the source is always given. With the exception of such quotations, this thesis is entirely my own work.
- I have acknowledged all main sources of help.
- Where the thesis is based on work done by myself jointly with others, I have made clear exactly what was done by others and what I have contributed myself.

Signed:

---

Date:

---



University College London

## *Abstract*

Faculty Name

Department of Computer Science

Master of Science

### **Song plays: Using machine learning to model the temporal patterns of music listeners**

by Badrul ALOM

While areas of customer recommendation have received a lot of attention in recent years, predicting users behaviour, particularly the timing of their actions, has not had as much focus. The modeling of temporal processes representing human behaviour can be applied across many industries. In this research we examine different methods for predicting a user generated event based on analyzing the timing of their past interactions. Our context is that of music listening and we examine a Bayesian inference model, logistic regression, a linear and RBF SVM, and a RNN-LSTM model. We find that precision and recall pose a trade-off, with a linear SVM being better at the former and an RBF model being better at the latter. We also find that an RNN model achieves relatively good result with a low amount of training data and expect that as computational power improves it will lead to top in class results in this area.



## Chapter 1

# Introduction

The modeling of events is useful across industries. For instance the times at which a customer makes an online purchase can help determine the optimal periods for target marketing. The times at which public transport users tend to travel can help better manage resources to meet demand. The times at which a medical illness re-occurs can help predict future episodes.

In all these cases modeling the temporal behaviour of the system is important in predicting the occurrence of the next event. Within the field of recommender systems, the area of predicting *what* a user would be interested in has received extensive attention in recent years, but *when* they would be interested in it, less so. In this research we look at how we can model the temporal behaviour of users, in order to help gauge their interest in an event at current time  $t$ , conditional on their history,  $h$ . More formally this is known as a temporal point process and we examine the existing methods for modeling such processes as well as recent innovative experiments with applying deep learning to the problem.

The research indicates that achieving a high precision and recall score on temporal point process problems requires a highly non-linear model, and the issue of whether deep learning can provide breakthroughs in this area is explored. We compare various classical linear and non-linear techniques with an RNN-LSTM model, and find our results on the applicability of deep learning to be inconclusive though positive enough to indicate the need for further research in the area.

## 1.1 Context

We take as our context for this research, the goal of estimating the probability that a user of a home-audio device would like to listen to music at a time period  $t$ , given their play history  $h$ . One application of this research would be to allow home audio devices to recommend music to a user at an opportune time. It could then also be extended for other activities.

The goal will be evaluate the effectiveness of several different machine learning methods. The research was guided by Emotech Ltd., a home audio hardware and software company and the creators of Olly [14].

## 1.2 Data

The dataset being used in this analysis is the LastFM1k dataset, which is freely available online and contains the listening history of a thousand LastFM listeners. It consists of a series of timestamps denoting when user started played a song. We wish to learn the temporal patterns of a user's behaviour in order to predict the next item in the sequence - a play or non-play event.

The dataset contains the timestamp, user ID, and track ID of users listening habits over a number of years (2005-2009).

### **1.3 Structure of the report**

We first perform a literature review of



## Chapter 2

# Literature Review

Analyzing event data presents its own unique challenges and questions. What are the ways to represent the data? What types of stochastic models are appropriate for explaining the structure in the data? How can we measure how well the data is described by a particular model?

Within literature the problem is referred to as event prediction, sequence prediction, or point process modeling. They represent event data in one of three ways: as event counts, as event times, and as inter-event intervals (fig. 2.1).

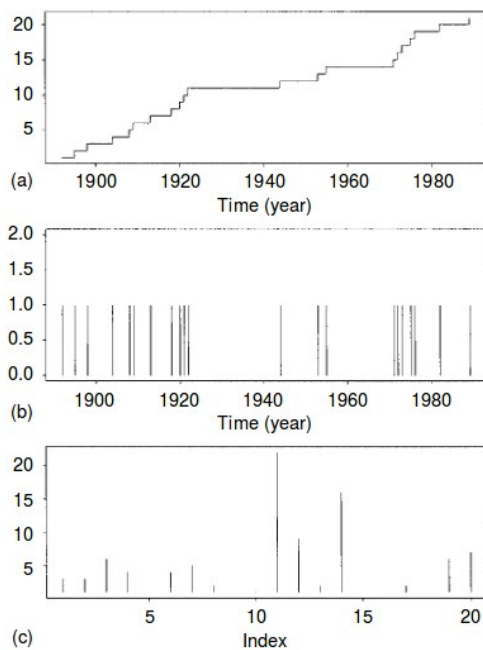


FIGURE 2.1: Three different representations of the same point-process  
a) event count b) date of occurrence c) interval time between floods

Of these inter-event time is the most common form of representing the data, for example in the times between financial transactions [7] while an example of the counting process is the scoring of goals in soccer [9].

## 2.1 Poisson Point Process

Poisson processes [10] are characterized by the fact that the probability of an event in a small time interval is independent of the event at all other times, and their probability or conditional intensity function follows a Poisson distribution. More formally we describe a Poisson point process,  $\xi$  as:

$$\Pr[\xi(T_i) = k_i, 1 \leq i \leq n] = \prod_i e^{-\lambda \|T_i\|} \frac{(\lambda \|T_i\|)^{k_i}}{k_i!}.$$

where  $\xi(T_1), \dots, \xi(T_n)$  are independent for separate subsets of  $T_1, \dots, T_n$

## 2.2 Temporal Point Processes

A temporal point process is the modeling of event data as a sequence of a fixed period intervals. It can be characterized in three ways: as an ordered list of event times  $T_1, \dots, T_n$ , as inter-event times  $U_1, \dots, U_n$  where  $u + i = T_i - T_{i-1}$ , or as a counting process where  $N(t)$  is a count of the number of events occurring before time  $t$  [1].

In the case of music listening we can model a sequence of events and non-events, where  $t_i$  can either be 0 (did not play music) or 1 (played music). This is the form we use in all but the Bayesian model as described in the next chapter.

### 2.2.1 Conditional Intensity Function in Point Processes

The conditional intensity function,  $\lambda(t)$ , represents the infinitesimal rate at which events are expected to occur around a particular time  $t$ , conditional on the prior history of the point process prior to time  $t$ .  $\lambda$  may be estimated parametrically [6] or via a parametric model. With the Poisson process,  $\lambda$  depends only on  $t$ . In other point processes it also depends on the history preceding  $t$ .

If this conditional intensity remains constant over time it is referred to as a homogenous or stationary point process [5]. If however it can vary with time it is known as an inhomogenous point process, such as the self-exciting (Hawkes) process where the intensity is determined by previous events through the parametric form  $\lambda(t) = \mu + \beta \sum_{t_i < t} g(t - t_i)$  and where  $g$  is some non-negative kernel function.

The logit function used in generalized linear model, can also be thought of as a conditional intensity function and has been used to develop sophisticated models such as [2] and [18].

However as noted by Wass et. al [22], point process models using a conditional intensity function often make various parametric assumptions about the latent dynamics governing the generation of the observed point patterns. As a consequence, model misspecification can cause significantly degraded performance using point process models.

## 2.3 Deep Learning Models

In recent years deep learning has demonstrated the power to learn hierarchical non-linear patterns on large-scale datasets [12] through multiple layers of abstraction (e.g. multi-layer feedforward neural networks). It has achieved state-of-the-art performances on a wide range of applications, such as computer vision [11], natural language processing [20], and protein structure prediction [13].

However it has not been applied to temporal point processes until recently with Xiao et. al [22] applying Generative Adversarial Networks (GANs) to the problem. GANs consist of two neural network models - a generator tasked with generating (i.e. predicting) a future sequence of events based on the history, and a discriminator tasked with detecting the true (ground truth) sequence amongst the generated ones.

For measuring the loss between a generated and true sequence, the authors found the Wassertein-Distance [15] performed better than Maximum Likelihood Estimate (MLE) which they remarked "may suffer from mode dropping or get stuck in an inferior local minimum".

Their findings showed that where as parametric point process models work better with problems where a parametetric form exists, with real world data a GAN model with Wasserterin-Distance outperformed all other models.

### 2.3.1 Recurrent Neural Networks (RNN)

Within this research we look at RNN models. RNNs are a type of artificial neural network designed to recognize patterns in sequences of data, such as text, genomes, handwriting, the spoken word, or numerical times series data emanating from sensors, stock markets and government agencies. Whilst a traditional Feed-Forward network [19] has input nodes, hidden layers, and an output layer, with data flowing in one direction only, RNNs allow for the hidden state from one timestep of the neural net to be an input into the next (see fig. 2.2).




FIGURE 2.2: RNN with 3 timesteps

RNN models can employ different methods for the propogation of the hidden state over time. One such well known method is Long Short-Term memory (LSTM) [17]. Here the hidden state is the product of a further four layers that interact in a way as to learn what information to retain and what information to throw away.




FIGURE 2.3: LSTM

Finally, as this report was being finalized, very recent research [23] came to light that is worth mentioning. In this the authors view the conditional intensity function of a point process as a non-linear mapping between the predicted transient occurrence intensity of events with different types, and the model input information of event participators, event profile and the system history.

They then go onto model this complex nonlinear mapping using recurrent neural networks, with logistic regression as a baseline model. Most interesting in their approach of utilizing two RNNs: one for modeling the time-series data and associated features (similar to the approach in our research) and a second to model long-range dependency over history with arbitrary time intervals; specifically a sequence of event type and the time since the previous event. In this way they can capture both the temporal based patterns, as well as the non-temporal, event-correlation patterns; which approximates to a conditional intensity function based on inter-event time.

The favoured evaluation metrics in the research were are precision, recall, f1 score, and confusion matrix.

## **2.4 Summary**

While not an exhaustive review of techniques we have seen some of the ways in which point processes can be modelled using traditional conditional intensity based models, as well as more recent areas of experimentation using deep learning. In the next chapter we layout the approach and methods we wish to explore in the context of this research.

## Chapter 3

# Experimental Design

In this chapter we describe the different methods that were assessed for the task of predicting whether a user is interested in listening to music at time  $t$  given their play history  $h$ . The methods, in order of gradually increasing sophistication, are as follows:

1. Baseline model
2. Bayesian Inference
3. Binary Logistic Regression
4. Linear SVM Classifier
5. Non-Linear SVM Classifier
6. Recurrent Neural Networks

All methods, except for the Bayesian Inference method, require the data to be structured as a time-series. The Bayesian method adopts a different approach that requires the data to be aggregated into half-hourly buckets of a week.

The data preparation that was performed is described first in the next section, followed by an explanation each method, and the evaluation criteria for assessing the methods.

### 3.1 Data preparation

#### 3.1.1 Data transformation

The analysis was carried out in Python (via Jupyter notebooks) running on Ubuntu. The raw data consisted of timestamps of when a song was played and a user ID. These were loaded as-is into a SQLite3 database in order to reduce the need to repeat data preparation steps. The methods themselves utilized Scikit-learn for all models, save for the RNN model which used Tensorflow.

UserIDs were converted to integer (e.g. 'User0005' became '5') and a period lookup table was created at  $n$  minute intervals, to which all timestamps in our main dataset could be mapped to.  $n$  was chosen to be 30 although it is possible to re-run the analysis for other levels of granularity.

More significantly the data, which contained entries for the times at which each user listened to music, was supplemented with all the times they did *not* listen to music, between their date of their first and last play. This was required in order to generate a sequence of play and non-play events. It was noted however that in cases where a user stops listening to music for a long period, then listens to it sporadically, we would get an even greater imbalance of data (see later).

### 3.1.2 Time-series fetures

The approach to feature selection for this research was to devise an initial of features based on the preliminary analysis, and use that list across all models.

The time series features that were derived from the raw data were time-lags:  $t, t-1, t-2, t-3, t-4, t-5, t-12hrs, t-23.5hrs, t-24hrs, t-24.5hrs, t-1wk, t-2wks, t-3wks, t-4wks$  where  $t$  is the event being predicted.

$t-1$  to  $t-5$  represent user activity in the previous 2.5 hours. The remaining time-lags were chosen to represent half-day, daily, and weekly cycles, with additional empahsis around the -24 hour mark due to the daily patterns observed in the preliminary analysis (see next chapter).

In addition a few non-time lag features were selected. One of which was the number of hours away from 5pm in either direction (so a timestamp at 4pm and 6pm would both equal 1), based on the observations in the preliminary analysis of 5pm being a peak hour. The others were binary features representing the day of the week (isMon, isTue etc.). Again these were based on the patterns observed in the preliminary analysis.

To improve the speed and quality of our models, rows where all time-lags indicated a non-play were removed from the dataset. Only 0.62% of such cases were found to be a play-event at time  $t$ .

## 3.2 Validation and Test dataset selection

Our working dataset was a subset of the full 1000 user dataset, and comprised of 4,217,228 rows of training data across 97 users. Of this a random sample of 100,000 rows was taken on which 5-fold cross-validation.

Due to the time it takes for training cross validation was not performed on the RNN model, instead 10 randomly selected users were held back from the main training dataset and a random sample of 10,000 rows were taken from each user.

### Data Imbalance

Data imbalance is when the training data is when one or more of the classes is under-represented in the dataset. Of the 4,217,228 rows in our data set, 361,081 (8.6%) were play events and 91.4% were non-play events.

This could lead to models that achieve a high accuracy score by following either or both of the following two heuristics: \* Predict non-event for everything \* Predict  $t$  will be the same as  $t-1$

There are several methods for dealing with data imbalance [3] including restricting input data, having a weighted loss function, and using recall as an evaluation measure.

In this research class weights were employed to automatically adjust weights inversely proportional to class frequencies in the input data  $n_{samples}/(n_{classes} * np.bincount(y))$ .

This would have the impact of encourgaing the model to predict play events more frequently. However our goal is to encourage the model to predict play events, while minimizing the number of false postivies on the predictions (i.e. we attach less cost to false-positives for non-event prediction). Encouraging this sort of behaviour requires not only class weighting, but also sample weighting to give a higher penalty to false positives for play-predictions. This strategyh was also employed in the models as shown in the following Python snippet (1 represents a Play event):

$$sampleWeights = 1 + (y[:,] == 1) * (x[:, 1] == 0)$$

### 3.3 Methods

#### 3.3.1 Baseline Model

Our baseline model will be to assume  $t = t - 1$ , that is to say a person will listen to music at period  $t$  if, and only if, they listened to music in the period immediately prior. As music listening events tend to be clustered (people listen to music in batches) the accuracy of the baseline model is expected to be fairly high. However it will not be able to predict the first play of a listening session, or where the listening session duration lasts no longer than a single period, which in the experiments is defined as a 30 minute period.

#### 3.3.2 Bayesian Inference

We employ a simple Bayesian Inference approach by utilizing the Beta-Binomial model. The Beta-Binomial model is built upon the weekly patterns observed in the preliminary analysis. Conceptually it seeks to build up a users personalized weekly listening profile as observations come in, using a Beta-Binomial probability distribution, with the priors based on the population as a whole. We then assess how effective this profile is at predicting listening events for that user.

Prior probabilities were calculated for each half-hourly period in a week ( $24*2*7=336$  timeslots). Fig 3.1 shows the calculations for the first 2.5 hours of a Sunday (d-hour-hh format).

The likelihood function which represents the probability of a user listening to music was defined as a binomial distribution, where  $k$  is the number of plays in a given period,  $n$  is the sum of plays and non-plays, and  $\theta$  is the unknown probability parameter for the binomial distribution.

$$\binom{n}{k} p^k (1 - p)^{n-k}$$

As the Beta distribution is a conjugate prior to the Binomial the model can be reduced to:  $Beta(\alpha + P, \beta + Q)$  where  $P$  is the count of plays and  $Q$  is the count of non-plays.

Our parameters for the beta distribution,  $\alpha$  and  $\beta$ , are derived from the training set, with an estimate of the mean for each half-hourly time period as shown in fig 3.1.

Timeslot	mean	var	a	b
1-00-1	0.098577	0.010807	0.711953	6.510370
1-00-2	0.092327	0.011242	0.595911	5.858451
1-01-1	0.090256	0.011784	0.538632	5.429205
1-01-2	0.089523	0.011741	0.531937	5.409996
1-02-1	0.087637	0.011772	0.507577	5.284259

FIGURE 3.1

We first calculate the probability of a play (total plays in period / count of plays and non-plays) per user, then take the mean and variance across users.  $a$  and  $b$  are then determined as:  $a = (\frac{1-\mu}{\sigma} - \frac{1}{\mu})\mu^2$  and  $\beta = \alpha (\frac{1}{\mu} - 1)$ .

Finally we convert the probabilities into a binary outcome by optimizing for a threshold  $\lambda$  at which we predict a play event.

### 3.3.3 Binary Logistic Regression

This method (and the subsequent methods) adopts a more classical time-series approach to the prediction problem by constructing a dataset as a sequence of play events at time  $t$   $Y_t = 1$  and non-play events  $T_t = 0$  with  $t$  representing a time period of a fixed interval, which for our research will be 30 minute chunks. We seek to calculate the probability of an event in the current time period  $t$ , given the history of events:  $p(Y_t = 1 | Y_h)$ , for each individual user together with additional non-time-series features (see .

Our binary logistic model is therefore defined as:

$$p(Y_t = 1 | Y_h) = \sigma(w^T x + b)$$

$$p(Y_t = 1 | Y_h) = 1 - p(Y_t = 0 | Y_h)$$

with  $\sigma$  being the sigmoid function defined as:

$$\sigma(x) = \frac{e^x}{1 + e^x}$$

Determining the optimal weights and constant can be determined by maximization of the log-likelihood or the minimization of the negative log-likelihood [**NegLog**]. The latter is given as:

$$C \sum_{i=1}^n \log(\exp(-y_i(X_i^T w + b)) + 1)$$

for class  $C = 1$ .

In addition we will be comparing results with and without the L2 regularizer term:  $+1/2w^T w$

Regularization is technique for encouraging the model to converge on a smaller set of strong features rather than a wide set of weak features and is important in helping us determine which time-lags and additional features are important and which are not.

### 3.3.4 Linear SVM Classifier

SVM models seek to determine a separation plane between classes based on the support vectors, the data points closest to the decision boundary.

A linear SVM regression model performs this through the Epsilon Intensive loss function [21]. The objective becomes to *minimize*:

$$\max(0, \|(y_i - w_i x_i - b) - \epsilon\|)$$

In other words we ignore cost functions that are within a certain margin  $\epsilon$ . In our case this may be of importance in cases where the probability of user listening to music is close to the decision boundary, which may be the case for the very first song played at the start of a session.



### 3.3.5 Non-Linear SVM Classifier

Here we extend a regression model to include a Gaussian RBF kernel. A Gaussian kernel is a popular method for modeling non-linear decision boundaries, so for example if the decision to play music is based on some non-linear combination of day of the week, time of the day, and time since last play. Our model becomes:

$$p(E = 1) = b + \sum_{i=1}^N w_i RBF(x, x_i)$$

where the RBF kernel is defined as:

$$K(\mathbf{x}, \mathbf{x}') = \exp\left(-\frac{\|\mathbf{x} - \mathbf{x}'\|^2}{2\sigma^2}\right)$$

Note that our actual implementation in Tensorflow makes use of the more computationally efficient algorithm as defined by McClure [16]. This restated method has a parameter  $\gamma$  that takes the place of  $2\sigma^2$  and which requires tuning.

Mini-batching was also performed with a batch size of 2000 and 5 iterations of the dataset, as it was found that rows of more than 5000 caused memory errors, and iterations beyond 5 showed minimal or no decrease in loss).

### 3.3.6 Recurrent Neural Networks

The construction of the RNN model requires a number of hyper-paramters that we will discuss in turn. They are:

spacing

- steps: How many time steps to use (= batch depth)
- learning rate: Learning rate of backprop algorithm
- Hidden units: Number of units per hidden layer
- layers: Number of hidden layers
- class1Weighting: Weighting to apply in the cost function for labels that are a play event
- user iteration: How many random users to select for training)
- samplesPerUser: How many mini-batches to select for each user
- batchRows: How many random periods to select in each mini-batch (=batch height)
- batch iterations: How many iterations to perform on one batch

#### Batch shape

We provide the RNN model with a sequence of historical data at times  $t - 1..t - n$  with  $n > 1$ ). This is fed into the RNN as separate time steps in forward temporal order (earliest time-steps are processed first). The hidden state would propagate information forward at each time step, until it was used to predict the outcome at time  $t$ .

Practically speaking this meant feeding in the data with input shape (batch rows, time-steps, features). Some points of interest to fellow researchers are:

1. The features dimension here does not need to contain all time-lags as we do in our other time-series models. However it must contain  $t-1$ .
2. The time-steps are 'unrolled' within Tensorflow and fed into the LSTM. In this research we experiment with different lengths of time-steps, with 48 (half-hour periods ) representing a day, 336 for a week, and 672 for 2 weeks.
3. When the data is unrolled, time step  $t$  for all rows in the batches are processed together as one block, before moving onto the next time-step.
4. Constructing the 3-d shape often requires building them up in slices. A significant speed up was observed in Python when using a pre-allocated array vs. appending to it.

### Samples and iterations

A challenge in the research was to balance the number of data samples taken from each user, with the number of users to sample from, and the number of iterations to do on each sample. As we see in our results this can lead to long training times. A distinction between batch rows and samples per user was required for computational reasons to reduce the demand on RAM per mini-batch, vs. feeding in one large mini-batch per user.

### Cost function and weighting

A weighted soft max cross entropy loss function was used to determine the cost (called in Tensorflow through the `tf.losses.sparse_softmax_cross_entropy` function).

A fully worked out example of this is shown in fig 3.2. It highlights the importance of the weighting that was applied to class 1 (play events) and why choosing the right weight is crucial to prevent the model drifting to predicting all 0's over time.

Logits		Softmax			Natural Log		Pick target	Weight	Cost	Actual label
0	1	0	1	Sum	0	1	0 or 1			
0.501	0.316	0.55	0.45	1	-0.60	-0.79	0.60	1	0.605	0
0.527	0.265	0.57	0.43	1	-0.57	-0.83	0.57	1	0.571	0
1.998	-1.158	0.96	0.04	1	-0.04	-3.20	0.04	1	0.042	0
0.501	0.316	0.55	0.45	1	-0.60	-0.79	0.60	1	0.605	0
0.500	0.316	0.55	0.45	1	-0.61	-0.79	0.61	1	0.606	0
1.818	-0.959	0.94	0.06	1	-0.06	-2.84	0.06	1	0.060	0
0.501	0.316	0.55	0.45	1	-0.60	-0.79	0.60	1	0.605	0
0.501	0.316	0.55	0.45	1	-0.60	-0.79	0.60	1	0.605	0
0.501	0.316	0.55	0.45	1	-0.60	-0.79	0.60	1	0.605	0
0.501	0.316	0.55	0.45	1	-0.60	-0.79	0.60	1	0.605	0
0.501	0.316	0.55	0.45	1	-0.60	-0.79	0.60	1	0.605	0
0.677	0.112	0.64	0.36	1	-0.45	-1.02	0.45	1	0.450	0
0.554	0.263	0.57	0.43	1	-0.56	-0.85	0.85	6	5.096	1
0.501	0.316	0.55	0.45	1	-0.60	-0.79	0.60	1	0.605	0
1.768	-0.920	0.94	0.06	1	-0.07	-2.75	0.07	1	0.066	0
0.494	0.317	0.54	0.46	1	-0.61	-0.79	0.61	1	0.609	0
1.584	-0.737	0.91	0.09	1	-0.09	-2.41	0.09	1	0.094	0
1.448	-0.623	0.89	0.11	1	-0.12	-2.19	0.12	1	0.119	0
0.501	0.316	0.55	0.45	1	-0.60	-0.79	0.60	1	0.605	0
0.521	0.285	0.56	0.44	1	-0.58	-0.82	0.58	1	0.582	0
0.501	0.316	0.55	0.45	1	-0.60	-0.79	0.60	1	0.605	0
Av. cost		0.475		Av. cost	0.687					

FIGURE 3.2: Softmax cross-entropy loss

The highlighted row is the only play event in the batch, which is not an atypical scenario. The logits predict this row to be a non-play event. Without weighting, our cost for this row would be 0.85, but after a class 1 weighting multiplier of 6 is applied it becomes 5.096.

As we see in the next chapter, the weighting play a crucial role in getting good results, and too low or too high a weighting leads to predictions drifting in favour of no-plays to all-plays.

### 3.4 Evaluation criteria

Deciding on an appropriate evaluation measure requires careful consideration to the costs attached to different predictions. In our case the cost of suggesting music when the user is likely not interested is higher than not playing music when they likely not interested.

We also wish to rate models higher if they manage to predict the first play of a listening session, as this is seen be harder than simply predicting a play event, when the last  $n$  events were also a play event.

With these in mind a number of possible evaluation criteria were looked at. The most-straight forward of which was *accuracy*. This computes the count of correct predictions as a fraction of the total number of predictions. While this is an intuitive measure, it would not distinguish between models that had a high accuracy on the play-events vs. a high accuracy on non-play events.

To do this we look at precision. Precision (P) is defined as the number of true positives over the number of true positives plus the number of false positives. A positive in this case is a play-event so our precision equation becomes:

$$Precision = \frac{CorrectPlayPredictions}{TotalPlayPredictions}$$

This will be measuring our models on how many of their play event guesses were correct which takes care of one of our considerations outlined above. However it does not distinguish between models that predict a play event only when the previous event was a play (a 'safe bet' and so a high precision score) vs. models that predict correctly in cases where the previous event was not a play and are therefore harder to get right. We turn to another measure to aid with this: recall.

Recall is defined as the number of true positives over the number of true positives plus the number of false negatives. A false negative is where a model predicts a non-event, when in fact there was an event. For example if we predicted 100 plays correctly but there were in fact 110 plays in the dataset, then recall would be  $100/(100 + 10) = 91$

$$Recall = \frac{CorrectPlayPredictions}{TotalPlayPredictionsInDataset}$$

Precision and recall in combination was therefore our chosen evaluation method. In addition tests were carried out examine how well the models performed on predicting the first play event in a series.

### 3.5 Summary

We have described how our data was transformed to make it useful for our experiments. In particular how we had to convert our list of play events into a list of play and non-play events for every period. By doing this we are faced with an imbalance of data as non-play events make up 91% of the data and so we employ a weighting strategy to counteract this.

We then described the methods we will be utilizing in our experiments, consisting of a Baseline model which is simply to assume  $t = t - 1$ , a Bayesian model which build up a weekly profile of listening habits for each user, Logistic and SVM models that apply classical machine learning time-series prediction techniques to our problem, and finally a deep learning model in the form of an RNN-LSTM where we seek

to utilize the capabilities for it to learn temporal patterns through a hidden memory state.

Finally we looked at different ways for measuring the success of our problem, with precision and recall being the preferred choice.

In the next chapter we shall present the results of our experiments and a discussion of what this tells us about learning user-level temporal patterns and performing event prediction.

## Chapter 4

# Results

We begin our discussion of the results with preliminary analysis of the data that helped shape the experimental design. Here we seek to understand what are some of the overarching patterns in music listening habits and how these look at an individual level.

After this we present a summary of our results followed by a discussion of the performance of each individual method.

### 4.1 Preliminary analysis

#### 4.1.1 Daily play patterns

By grouping track plays into 30 min intervals and aggregating by periods within a day, we see a clear daily pattern with music listening hitting a peak at around 5pm and a trough at around 6am.

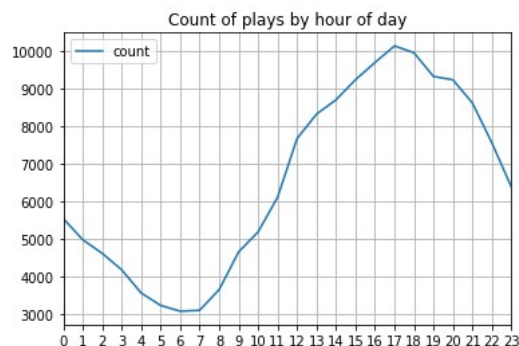


FIGURE 4.1: 5-5.30pm is peak listening time

Zooming out to view the pattern across an entire week in figure 4.2, we see that the daily pattern occurs across every day of the week with weekends having a lower total number of plays.

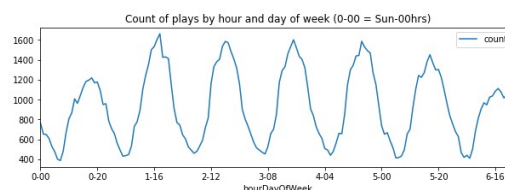


FIGURE 4.2: Most popular times to listen to music across all users

At a high level therefore one can get good accuracy by simply anticipating music demand to peak at 5pm. However if we select two users at random, we see (see

fig. 4.3) that these daily patters are not as strongly discernable. This demonstrates why models modeling the high levels patterns is not enough for individual user prediction.

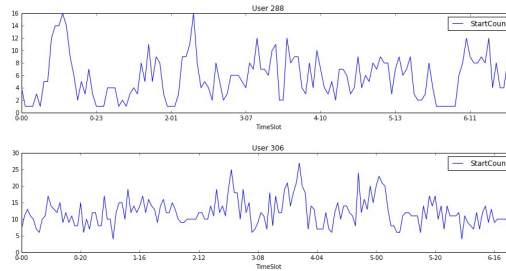


FIGURE 4.3: Most popular times to listen to music by individual user

### 4.1.2 Inter-event times

The dataset contains a timestamp associated with each user. This does not necessarily mean the user played a song in its entirety. Analysis shows plenty of cases where the interval time between tracks was a few seconds suggesting the user skipped tracks.

Figure 4.6 shows a frequency plot of intervals. Intervals beyond 30 minutes continue the exponential decrease and are not shown. We see that while the mode is on par with a typical song length, there is a significant number of plays that lasted under 5 minutes.

### 4.1.3 Time-series analysis

Here we examine our data once it has been transformed in a binary sequence of events (1) and non-events(0). We seek to understand better how an optimizer may perform based on traits of the data.

We begin with assessing how well our baseline model may perform based on assuming  $t = t - 1$ . Fig 4.4 shows that the 76% of Plays, also had a play in t-1. However simply using this as a rule would also capture 2.2% of non-plays.

	Count	%
<b>Plays</b>	361,081	
of which t-1 is a play	274,985	76.2%
of which t-1 is a non-play	86,096	23.8%
<b>Non-Plays</b>	3,856,147	
of which t-1 is a play	86,088	2.2%
of which t-1 is a non-play	3,770,059	97.8%
<b>Total</b>	<b>4,217,228</b>	

FIGURE 4.4

Furthermore the 23.8% of Plays that did not have a Play in the prior period are harder to predict yet of more interest as they represent the beginning of the listening period and therefore more useful to a music recommender system.

Given the daily patterns we have seen, it might be reasonable to assume that looking at the same period 24 hours prior may be a good indicator of whether t is a play event. However as we see from fig. 4.5 this is not a reliable indicator either.

	Count	%
<b>Plays</b>	361,081	
of which t-1 is a play	274,985	
of which t-24hrs is a play	102,522	37.3%
of which t-24hrs is a non-play	172,463	62.7%
of which t-1 is a non-play	86,096	
of which t-24hrs is a play	23,478	27.3%
of which t-24hrs is a non-play	62,618	72.7%

FIGURE 4.5

What both of these results tell us is that fairly high precision score of around 76% ought to be possible purely based on t-1 but going above this while having a good precision score on the non-play events will be harder.

Finally it was found that of rows where all timelags had a non-play event, 0.62% of rows had a play event. As this was a low number it was deemed safe to remove all such rows from the dataset in order to improve the speed and quality of analysis.

#### 4.1.4 Outliers

The data was checked for any unusual outliers that may impinge upon the goal of developing a model to predict user behaviour. An analysis of plays by user reveals a high amount of variance between users on how many tracks are played.

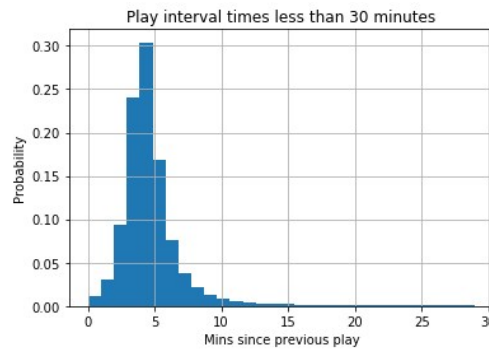


FIGURE 4.6

For our purposes these are include as evidence that the user was interested in playing music at time  $t$  and therefore treated as a Play event.

We can also assume that the song plays are not independent of one another, in that the probability of a play event at time  $t+1$  is significantly higher if there was an event at time  $t$ .

Further analysis showed one user in particular with very high amount of plays, with very low durations, suggesting it was likely to have been generated by a bot, possibly a LastFM test. This was excluded from the dataset.




FIGURE 4.7: Total play count by user

## 4.2 Main results

### 4.2.1 Summary

Fig. ?? shows the results from our experiments. We see how that the baseline model scores highly on precision but low on recall. In other words assuming  $t = t - 1$  is often quite often correct given that a group of songs are often played consecutively, but it fails to pick up the start and end points of a sequence of song plays, hence the lower recall.

Model	Train rows	Precision	(+/-)	Recall	(+/-)
Baseline	500k	79%	0.02%	13%	0.03%
Bayesian Inference	500k	41%	0.5%	13%	0.5%
Logistic Regression	500k	67%	0.5%	73%	0.1%
Linear SVM Classification	500k	81%	0.04%	68%	0.1%
SVM RBF Classification	100k	77%	0.01%	76%	0.1%
RNN-LSTM (T-1)	50k	70%	0.1%	69%	0.1%
RNN (T1 .. T5weeks)	50k	70%	0.8%	64.0%	0.7%

FIGURE 4.8: Summary of results. Note that both RNN models included non-time lag features (e.g. IsMon)

The Bayesian Inference model performed poorly on both measures. This would suggest that the difference between the population patterns and what is observed at an individual level differs too much for the prior probability to be accurate for any one individual user. However this does not rule out the Bayesian approach. An improvement might be to apply a logistic regression model in which the weights are continually updated with each new observation from the user. However such as an approach was not looked at in this research.

Our remaining models exhibit the tension between achieving a good accuracy and achieving a good recall. The Logistic regression model performs below average on precision but is second highest on recall. The Linear SVM and RMF regression models are on par with one another with the former achieving the highest score on prevision, and the latter on recall.

The Linear SVM model ignores probabilities that fall within a margin of the decision boundary during its optimization process. That this translates into performing well on precision may be due to placing lesser emphasis on ad-hoc play events that aren't indicative of a users general behaviour, as these probabilities would likely fall



near the decision boundary. Alternatively it could be better at dealing with the first play of a sequence whereby it often, but not always occurs at the same time each day; again leading to probabilities nearer the decision boundary.

The RBF model generally performs well on both metrics. Note that the RBF classification was restricted to 100k rows of training data as it was found the computational cost of 500k rows was too high so it had less training data to work with than the Linear SVM. The high recall rate of 68% indicates that inputs have a non-linear relationship to the output and that this is especially helpful to predicting the start and end of a play sequence as seen by the high recall score.

The RNN models were trained on 50k rows of data due to the computational cost of evaluating a large number of hyperparameters. Two types of models were tested - one with  $t - 1$  as the only time-lag information and one with all other time-lags included. Given the much smaller training set the results we observe are better than expected. However a strong caveat is that these results were hard to achieve and require a lot of hyper-parameter testing, which we discuss later in this chapter.

### 4.3 Beta-Binomial analysis

We can plot the prior probability for any given time period as shown in fig 4.9.




FIGURE 4.9

Here we see that the  $\theta$ , which represents the prior probability of listening to music in the timeslot, is most likely to be less than 0.1 for this timeslot. Note that this timeslot is for hour 17 (i.e. 5pm) which from our preliminary analysis we know is the peak listening time. The fact that the probability is so low suggests that users may often have weeks in which no music was played, and hence the prior probability for a timeslot is dragged down.

For our posterior function, the threshold at which we determine a Play event was determined by comparing the false positive rates with the true positive rates using a ROC curve (4.10). From this, 0.4 was selected as the optimal threshold at which to determine a Play event across all users.




FIGURE 4.10: ROC curve showing 0.4 as the optimal threshold

The results from the model are shown in figure 4.11. We see that the recall and precision of play events (as denoted by 1) is very low suggesting that relying on a Bayesian approach centered around a weekly profile of each users habits is not an effective method for predicting a play event for a new time period.

	precision	recall	f1-score	support
0	0.93	0.98	0.96	227823
1	0.40	0.13	0.20	19763
avg / total	0.89	0.91	0.89	247586

FIGURE 4.11: Beta-Bionmial Model Results

## 4.4 Logistic regression analysis

An advantage of logistic regression over non-linear models is that the coefficients are easily interpretable. While feature engoneering was not the focus of this research an examination of these can yield insight into which features may be unnecessary.

Field	Log Reg.
t-1	3.99677
t-2	0.66650
t-3	0.42272
t-4	0.40080
t-5	0.49478
t-23.5 hrs	1.12593
t-24hrs	0.55006
t-24.5 hrs	0.53480
HrsFrom5pm	-0.05715
isSun	-0.23220
isMon	-0.10579
isTue	-0.14689
isWed	-0.12977
isThu	-0.14111
isFri	-0.19522
isSat	-0.27923

FIGURE 4.12

Fig 4.12 shows that  $t - 1$  was by far the most important feature. This is expcted and it's importance crowds out the other features.

Interestingly t-23.5 is the second strongest time-lag and more significant than t-24hrs. This likely reflects the fact that having a consistent  $t - 1$  is more important than simply knowing what  $t$  was 24 hours prior.

The non-time lag features seem to have very little effect, which was also see in fig 4.13 where we do a forward stepwise regression are barely impacted by these features.

## 4.5 SVM analysis

Of the two SVM models, the RBF model was selected for more in-depth analysis due to its ability to better capture recall without too detrimental a loss in precision. An RBF kernel has two hyper-paramters that can significantly impact the results. The  $C$  parameter and the  $\gamma$ . Intuitively, the gamma parameter defines how far the influence of a single training example reaches, with low values meaning 'far' and




FIGURE 4.13: Stepwise regression, adding in fields from left to right

high values meaning ‘close’. The C parameter trades off misclassification of training examples against simplicity of the decision surface. A low C makes the decision surface smooth, while a high C aims at classifying all training examples correctly by giving the model freedom to select more samples as support vectors.

In order to determine optimal gamma and C values, a grid search was carried out with precision and recall of the test dataset plotted on a heatmap. In order to get through the many iterations of the data, the input data was reduced to 30,000 training samples. Fig 4.14 shows the final heatmap, after multiple iterations to home in on the optimal values, with precision on the left and recall on the right. The optimal values are deemed to be 0.6 for C and 0.76 for Gamma, although as the charts show there is some scope for movement around those values.




FIGURE 4.14: Heatmap of hyper-parameter search. Dark blue shade = better precision and recall

Finally, a note on the computational complexity of the RBF model. We used Scikit Learn’s implementation of the RBF kernel which is based on libsvm. The fit time complexity of this approach is more than quadratic with the number of samples, making it hard to scale to datasets with more than a couple of 10,000 samples [4]. While we managed to get as high as 100,000 samples in our training it was not considered a practical to go any higher thereby limiting the scalability of this approach and our ability to potentially obtain higher scores than what we already had.

## 4.6 RNN analysis

The RNN-LSTM model required a lot of hyper-parameter testing and as such the training dataset ended up being narrowed down to no more than 50k random samples and often as low as 20k. The eventual set of hyper-parameters that was found to work well *with early stopping* were:

spacing

- user iteration: 10 (How many random users to select for training)
- samplesPerUser = 20 # How many mini-batches to select for each user

- batchRows = 50 # How many random periods to select in each mini-batch (=batch height)
- batch iterations = 10 # How many iterations to perform on one batch
- steps = 672 # How many time steps to use (= batch depth)
- learning rate = 0.001 # Learning rate of backprop algorithm
- Hidden units = 250 # Number of units per hidden layer
- layers = 4 # Number of hidden layers
- class1Weighting = 8 # Weighting to apply in the cost function for labels that are a play event

This scored an average precision of 70% and an average recall of 66.5% over two separate runs, with 1 run taking approximately 10 hours to complete. The original number of user iterations was 20 but a consistent pattern was that if a model runs for too long, the results swing in favour of all plays, or no plays, depending on the weighting used. Periodic testing during training and early stopping when good performance is achieved is recommended.

Fig 4.15 shows the results of some of the other tests that were performed. The table has been shaded to aid understanding with darker shading what may be considered "better", whether this be more training data samples, a more complex model structure, or a better score.

Run number	1	2	3	10	12	13	13	15
Features list	t-1	t-1	t-1	t-1	t-1	t-1	t-1	t-1
User iteration	2	5	5	8	10	10	10	20
Samples per user	1000	1000	1000	50	50	20	5	20
Batch size	1000	1000	1000	50	50	30	30	50
Total input rows	2,000,000	5,000,000	5,000,000	20,000	25,000	6,000	1,500	20,000
Batch iterations	20	50	10	10	10	7	7	10
Time steps	336	672	672	672	336	672	672	672
Hidden Units	150	250	250	250	250	250	250	250
Layers	1	1	1	4	4	4	4	4
Learning rate	0.01	0.01	0.01	0.01	0.01	0.01	0.02	0.02
Imbalance data weighting	1	1	1	8	2	2	10	10
Precision	57.00%	50.00%	71.00%	70.00%	0.00%	0.00%	0.09%	0.09%
Recall	14.0%	23.0%	62.0%	69.00%	0.00%	0.00%	1.00%	1.00%

FIGURE 4.15: RNN results over multiple runs. Darker shading means better

Note that not all runs are shown - those that incorporated additional features beyond t-1 have been excluded.

### Hidden Layers

We see that the improvement in precision comes about most significantly through the use of 4 layers (see run no. 7) although a higher number of iterations of a mini-batch also seems to help improve precision (see run 2 vs. 3).

Yet too many iterations of a mini-batch adversely impacts recall. Good recall appears to be a product of a higher number of hidden units in a layer, a higher number of time steps being fed into the LSTM, and a higher number of layers. All these supports the idea that the ability to detect the start and end of a play session, requires a highly non-linear solution.

### 4.6.1 Declining scores

We also see a few results that stand out. They are runs 3, 12, and 13. Run 3 shows a reasonable score was achieved with just 1 hidden layer by performing fewer mini-batch iterations than run 2. While runs 12 and 13 showed a score of 0% by the end of their runs.

Both of these are attributed to an observation that the longer the model is trained the more likely it is to start predicting 'not a play event' for everything. We hypothesized that this is due to the data imbalance, which in over many training iterations leads the model to conclude getting the non-play events correct is more important than getting the play events correct.

We see this better in fig 4.16 which shows the precision and recall of every mini-batch sample. Note that in this chart, every 50 mini-batches represents the end of a random user, before another one is randomly selected. Fig. 4.16 shows us both the scores of each mini-batch training data (lighter colored markers), as well as the scores at the end of each random user, using randomly selected test data (darker symbols).




FIGURE 4.16: RNN results over multiple runs. Darker shading means better




FIGURE 4.17: RNN results over multiple runs. First 100 rows.

Note that it is generally easier to achieve a high precision score by only guessing play events infrequently, resulting in a lower recall score; and that in each mini-batch, less than 10% of rows would be a 'play event'.

From looking at fig. 4.16 it is hard to discern whether there are any patterns. For this reason we also include fig. 4.17 showing just the first 100 rows.

There are two insights we draw from these charts:

1. It is hard to tell if results are improving with each mini-batch however we do see in fig. 4.17 that there are fewer scores of 0% as training goes on.

2. After a certain point in training results fall off a cliff edge – the model predicts ‘No play’ for every event

Finally, an RNN-LSTM model has a computational complexity of per time-step and weight of  $O(1)$  [8]. While training still took some time, with better hardware applying the method to ever larger datasets is not intractable.

## Chapter 5

# Summary

### 5.1 Conclusion

Predicting the propensity of a user to listen to music at a certain time, based on their recent listening history can be applied to a range of other areas such as the propensity to purchase products, electricity usage, or the demands on a public transport system. The ability to accurately model these patterns is therefore of great significance to industry. Traditional temporal point processes modeling requires making assumptions about the data in order to build a prediction model and struggle with capturing highly non-linear patterns in a scalable manner.

Within this research we applied a range of techniques to the task of modeling temporal point processes, within the context of modelling the times at which users listen to music. Our evaluation measures were precision and recall.

As music listening is typically long periods of non-play events, followed by sequential periods of play events, the problem becomes that of a counter-balance between good precision and good recall. Improving recall requires predicting the start and end of a sequence while precision favours restricting predictions to when  $t - 1 = 1$ . This effect is seen in our baseline model which scored 79% on precision and 13% on recall after 5-fold cross validation.

It is also seen in the linear SVM models which is less impacted by cases that fall close to the decision boundary, as is the case with the start and end of a sequence of play events. From a practical perspective, predicting the start and end of series of events is what is of most interest and harder to do.

We found that while linear models are adept at a high precision score, they were poor at recall. Non-linear models performed much better on recall without compromising precision too much. An SVM model with an RBF kernel achieved a precision score of 77% and a recall score of 76%. However such a model is not deemed scalable to large datasets.

We also examined an RNN-LSTM model. This method suffers from requiring a large number of hyper-parameter tuning which results in a long training time. However once these hyper-parameters have been settled upon (such as through analysis of a subset of data), the method can be applied to a much larger set of data than is possible with an SVM RBF model, although it is noted that better hardware can make the difference between hours, days, or weeks of training. Our RNN-LSTM model scored an average precision score of 70% and 69% based on a small subset of data.

As our literature review at the start of the report showed, the application of deep learning to temporal point processes is still in its infancy but shows promising signs. Our own research is inconclusive though positive enough to suggest further research would be worthwhile. Some caution for application in industry is raised, due to the time required for hyper-parameter turning on large datasets. If there is a lot of

variance within a dataset, then the combination of multiple iterations and multiple hyper-parameter searches may make it an intractable solution. Nevertheless with processing power continually increasing there is the potential for deep learning to transform our approach to temporal point processes, and further investigation in this area is most welcome.

## 5.2 Future research

There are several directions fellow researchers could choose to take this problem.

1. Methods of efficient searches of hyper-parameter tuning of RNN temporal point process models
2. Confirmation of the hypothesis that through further training RNNs can achieve a very high score on temporal point process problems
3. Evaluation of more sophisticated RNN structures such as that described in our literature review [23].
4. Evaluation of advanced Bayesian models such as Bayesian Logistic regression
5. Evaluation of Gaussian Point processes as an alternative to RNN models
6. Investigation into how quickly models can learn to model the behaviour of new users



## **Appendix A**

# **Appendices**

### **A.1 A. Further detail of experiments**



# Bibliography

- [1] P. Andersen and O Børgan. *Statistical Models Based on Counting Processes*. Springer, 1993.
- [2] Adrian Baddeley et al. “Logistic regression for spatial Gibbs point processes”. In: *Biometrika* 101.2 (2014), pp. 377–392.
- [3] Jason Brownlee. *8 tactics to combat imbalanced classes in your machine learning dataset*. 2015. URL: <http://machinelearningmastery.com/tactics-to-combat-imbalanced-classes-in-your-machine-learning-dataset/>.
- [4] Chih-Chung Chang and Chih-Jen Lin. “LIBSVM: a library for support vector machines”. In: *ACM transactions on intelligent systems and technology (TIST)* 2.3 (2011), p. 27.
- [5] Sung Nok Chiu, Dietrich Stoyan, and Wilfrid S. Kendall. *Stochastic Geometry and Its Applications*. John Wiley and Sons, 2013.
- [6] Nan Du et al. “Time-sensitive recommendations from recurrent user activities”. In: *Advances in Neural Information Processing*. NIPS. 2015.
- [7] Robert Engle and Jeffrey R. Russell. “Autoregressive Conditional Duration: A New Model for Irregularly Spaced Transaction Data”. PhD thesis. University of California, 1996.
- [8] Felix A Gers, Jürgen Schmidhuber, and Fred Cummins. “Learning to forget: Continual prediction with LSTM”. In: (1999).
- [9] A. Heuer, C. Müller, and O. Rubner. “Soccer: Is scoring goals a predictable Poissonian process?” In: *EPL (Europhysics Letters)* 89 (Feb. 2010), p. 38007. DOI: [10.1209/0295-5075/89/38007](https://doi.org/10.1209/0295-5075/89/38007). arXiv: [1002.0797](https://arxiv.org/abs/1002.0797) [physics.data-an].
- [10] J. F. C. Kingman. *Poisson Processes*. Clarendon Press, 1992.
- [11] Alex Krizhevsky, Ilya Sutskever, and Geoffrey E. Hinton. “ImageNet Classification with Deep Convolutional Neural Networks”. In: *Advances in Neural Information Processing*. NIPS. 2012.
- [12] Yann LeCun, Yoshua Bengio, and Geoffrey Hinton. “Deep Learning”. In: *Nature* 521 (May 2015).
- [13] Pietro Di Lena, Ken Nagata, and Pierre Baldi. “Deep architectures for protein contact map prediction”. In: *Bioinformatics* 28 (19 Oct. 2012).
- [14] Emotech Ltd. *Your robot with personality*. URL: <https://www.heyolly.com/>.
- [15] Léon Bottou Martin Arjovsky Soumith Chintala. “Wasserstein GAN”. PhD thesis. Facebook AI Research, 2017.
- [16] Nick Mclure. *Tensorflow Cookbook*. 2016. URL: [https://github.com/nfmcclure/tensorflow-cookbook/blob/master/04\\_Support\\_Vector\\_Machines/04\\_Working\\_with\\_Kernels/04\\_svm\\_kernels.ipynb](https://github.com/nfmcclure/tensorflow-cookbook/blob/master/04_Support_Vector_Machines/04_Working_with_Kernels/04_svm_kernels.ipynb).

- [17] Christopher Olah. *Understanding LSTMs*. 2017. URL: <http://colah.github.io/posts/2015-08-Understanding-LSTMs/>.
- [18] Tuomas Rajala. “A note on Bayesian logistic regression for spatial exponential family Gibbs point processes”. In: *arXiv preprint arXiv:1411.0539* (2014).
- [19] D. E. Rumelhart, G. E. Hinton, and R. J. Williams. “Learning internal representations by error propagation”. In: *Parallel distributed processing: explorations in the microstructure of cognition*. Cambridge, MA: Association for Computational Linguistics, 1986, pp. 1631–1642.
- [20] Richard Socher et al. “Recursive Deep Models for Semantic Compositionality Over a Sentiment Treebank”. In: *Proceedings of the 2013 Conference on Empirical Methods in Natural Language Processing*. Seattle, WA: Association for Computational Linguistics, 2013, pp. 1631–1642.
- [21] Vladimir N. Vapnik. *The nature of statistical learning theory*. Springer-Verlag, 1995.
- [22] Shuai Xiao, Mehrdad Farajtabar, and Xiaojing Ye. “Wasserstein Learning of Deep Generative Point Process Models”. PhD thesis. Georgia Institute of Technology, 2017.
- [23] Shuai Xiao et al. “Modeling the Intensity Function of Point Process Via Recurrent Neural Networks.” In: *AAAI*. 2017, pp. 1597–1603.